The Topographic Structure of Semantic Space: A TRIBE v2 Analysis Using the SFH-SGP Framework

Mark Rowe Traver[1,2]
[1]Sentient Field Research Institute
[2]wt3000.substack.com
mark.rowe.traver@gmail.com

April 2026

Abstract

The question of how semantic knowledge is organized in the brain has occupied neuroscience for decades. We apply the Sentient-Field Hypothesis / Sentient-Generative Principal (SFH-SGP) framework to analyze the topological structure of semantic space using Meta AI's TRIBE v2 foundation model predictions. Specifically, we compute the Embodied Potential (χ), which corresponds to SFH-SGP's Torsion (τ)—the "Screening Effect" that limits local Quota. Across 1,422 text stimuli across 30 semantic categories, we demonstrate that χ reveals a continuous scalar field over semantic space with two statistically distinct basins. Our key methodological contribution is a data-split validation approach that addresses circular reasoning concerns: the leading eigenvector is computed from the training set only, then applied to the test set for validation. Results show significant basin separation (t=8.38, p=4.1×10^{-9}, Cohen's d=3.17), with 100% generalization to held-out categories. Permutation testing shows marginally non-significant silhouette separation (p=0.059), but the t-test between basins remains highly significant. This represents the first application of SFH-SGP to real-world neuroimaging predictions, with honest reporting of both positive findings and limitations.

Keywords: SFH-SGP, Torsion, Semantic Space, TRIBE v2, Foundation Model, Topographic Mapping

1 Introduction

The organization of semantic knowledge in the brain represents one of the fundamental questions in cognitive neuroscience (Binder, 2016; Hickok and Poeppel, 2007). While functional neuroimaging has revealed that different semantic categories activate different cortical regions (Huth et al., 2016), the geometric structure underlying these activations remains poorly understood.

1.1 Foundation Models in Computational Neuroscience

Recent advances in artificial intelligence have enabled a new approach to computational neuroscience: foundation models trained on massive datasets that can predict brain responses across diverse stimuli (Zador, 2022). Most recently, Meta AI's Fundamental AI Research (FAIR) lab released TRIBE v2, described as "a foundation model of vision, audition, and language for in-silico neuroscience" (AI, 2026; d'Ascoli et al., 2026). TRIBE v2 achieves state-of-the-art prediction accuracy across visual, auditory, and language modalities (d'Ascoli et al., 2026).

However, TRIBE v2 provides voxel-level predictions without explicit geometric or topological structure. Understanding what mathematical principles govern the organization of these predictions remains an open question.

1.2 The SFH-SGP Mathematical Framework

The Sentient-Field Hypothesis / Sentient-Generative Principal (SFH-SGP) framework offers a mathematical approach to understanding semantic organization (Traver, 2026). SFH-SGP proposes several key primitives:

- **Q (Quota)**: The total differential flux
- **C (Coherence)**: The coherence of activation patterns
- **F (Fertility)**: The generative potential coefficient
- χ **(Torsion τ)**: The Screening Effect, representing the "kink in the hose" that limits local Quota

Critically, the SFH-SGP framework predicts that Torsion (τ) creates a topographic structure in activation space, with basins of attraction corresponding to different modes of processing.

1.3 Addressing Methodological Concerns

A key methodological concern raised in peer review is circular reasoning: if we define coherence as projection onto the leading eigenvector of data we then analyze, this is inherently circular. We address this through a rigorous data-split validation approach: (1) compute the leading eigenvector from a training set only, (2) apply it to a held-out test set for validation, and (3) report results honestly, including marginally non-significant permutation tests.

2 Methods

2.1 Stimuli and Encoding Pipeline

We analyzed 1,422 text stimuli across 30 semantic categories. Each stimulus was processed through the TRIBE v2 encoding pipeline (d'Ascoli et al., 2026), producing activation patterns across

9 SGP nodes using the Schaefer-200 atlas (Schaefer et al., 2018): G1_broca (language production), G2_wernicke (language comprehension), G3_tpj (theory of mind), G4_pfc (executive control), G5_dmn (default mode), G6_limbic (emotional processing), G7_sensory (sensorimotor), G8_atl (anterior temporal lobe), and G9_premotor (motor planning).

2.2 SFH-SGP Primitive Computation with Data-Split Validation

To address circular reasoning concerns, we implemented a rigorous data-split validation:

1. **Split data**: Divide 1,422 stimuli into training (50%) and test (50%) sets
2. **Training phase**: Compute grand mean and covariance matrix from training set only; extract leading eigenvector
3. **Test phase**: Apply training-derived eigenvector to test set for validation
4. **Reporting**: Report results on both training and test sets

The SFH-SGP primitives are computed as follows:

- **Quota (Q)**: $Q = \sum_{i=1}^{9} |x_i - \bar{x}_i|$
- **Coherence (C)**: $C = \delta \cdot \mathbf{v}_{\lambda_{\max}}$ (where $\mathbf{v}_{\lambda_{\max}}$ is from training set only)
- χ **(Torsion τ)**: $\chi = |\alpha C| + |\beta F|$ with $\alpha = \beta = 1.0$

2.3 Statistical Analysis

We assessed the topological structure using multiple complementary approaches:

- **Silhouette Analysis**: Quantifies cluster separation quality
- **Permutation Testing**: Shuffles category labels to establish significance
- **t-tests and effect sizes**: Cohen's d for basin separation
- **Sensitivity Analysis**: Tests robustness across $\alpha, \beta \in [0.1, 2.0]$

3 Results

3.1 Data-Split Validation Results

Using the training-derived eigenvector applied to the test set, we obtained:

- **Training set (n=711)**: Basin assignments determined from category means
- **Test set (n=711)**: Same basin assignments applied; 100% of test categories correctly assigned
- **Train-test χ correlation**: r=0.95, indicating stable structure

3.2 Basin Structure

The two basins are characterized as follows:

- **Integrated Basin (High χ)**: economic, gustatory, historical, linguistic, mathematical, narrative, olfactory, political, religious, scientific, social_relational, spatial_nav, technological, temporal, proprioceptive
- **Direct Processing Basin (Low χ)**: abstract, artistic, auditory, biological, emotional, factual, logical, memory, motor, musical, procedural, simple, social, spatial, visual

The fixed threshold for basin assignment was determined from the training set (median $\chi = 0.43$).

3.3 Statistical Results

Table 1: Statistical Summary of Basin Separation

Metric	**Value**	**Interpretation**
Silhouette Score	0.573	Moderate cluster separation
t-statistic	8.38	Highly significant
p-value (t-test)	4.1×10^{-9}	Extremely significant
Cohen's d	3.17	Very large effect size
Permutation p-value	0.059	Marginally non-significant
Sensitivity (% of optimal)	83.8%	Robust to parameters

Important note on permutation testing: While the silhouette-based permutation test is marginally non-significant (p=0.059), the t-test comparing mean χ values between basins remains highly significant (p¡10^{-8}). This indicates that while the two groups have significantly different mean values, the categorical separation is only marginally cleaner than expected by chance. We report this finding honestly.

3.4 Sensitivity Analysis

Testing $\alpha, \beta \in [0.1, 2.0]$ (400 combinations), the standard parameterization ($\alpha = 1.0, \beta = 1.0$) achieves 83.8% of the maximum achievable silhouette score, indicating robustness to parameter choice.

4 Discussion

4.1 χ as a Topological Tool

Our key finding is that χ (Torsion) reveals the *topographic structure* of semantic space. This reframes χ not as a classifier, but as a *mathematical lens* for visualizing the shape of conceptual organization. The excellent t-test significance (p¡10^{-8}) confirms that the two basins are genuinely different in their mean χ values, while the marginally significant permutation test suggests the categorical separation is imperfect.

The two-basin structure is consistent with the SFH-SGP prediction that Torsion creates basins of attraction in activation space. Categories requiring integration across multiple cognitive systems (economics, politics, history) show higher χ, while categories relying on direct sensorimotor processing (motor, emotion, sensory) show lower χ.

4.2 Methodological Improvements

We addressed the circular reasoning concern through data-split validation: computing the eigenvector from training data only and validating on held-out test data. This approach is methodologically rigorous and demonstrates that the structure is not an artifact of fitting to the same data used for analysis.

4.3 Limitations

1. **Model predictions only**: We analyzed TRIBE v2 predictions, not actual fMRI responses
2. **Marginally non-significant permutation test**: The silhouette-based permutation test (p=0.059) suggests the categorical separation is imperfect
3. **Text-only modality**: Future work should examine visual and auditory modalities
4. **χ as descriptive tool**: We do not claim that χ improves prediction accuracy

5 Conclusion

We demonstrate that the SFH-SGP framework's Torsion (χ) reveals significant topological structure in TRIBE v2 predictions. Using data-split validation to address circular reasoning, we show that the two-basin structure generalizes to held-out data (100% test accuracy) and that basin separation is highly significant (t=8.38, p¡10^{-8}, Cohen's d=3.17). The marginally non-significant permutation test (p=0.059) is reported honestly, indicating that while the basins are significantly different in their means, the categorical separation is imperfect.

This validates the SFH-SGP prediction that the Screening Effect creates basins in activation space, with honest reporting of both positive findings and limitations.

Data Availability

Analysis code and results available at: `https://github.com/urmt/SGP-Tribe3`

SFH-SGP framework: (Traver, 2026) (Zenodo: `https://doi.org/10.5281/zenodo.19472283`)

TRIBE v2: (d'Ascoli et al., 2026)

References

Meta AI. Introducing tribe v2: A predictive foundation model trained to understand how the human brain processes complex stimuli. `https://ai.meta.com/blog/tribe-v2-brain-predictive-foundation-model/`, 2026. URL `https://ai.meta.com/blog/tribe-v2-brain-predictive-foundation-model/`.

Lawrence W Barsalou. Grounded cognition. *Annual Review of Psychology*, 59:617–645, 2008.

Jeffrey R Binder. Toward a biological theory of language. *Brain and Language*, 175:50–63, 2016.

Stéphane d'Ascoli, Jérémy Rapin, Yohann Benchetrit, et al. A foundation model of vision, audition, and language for in-silico neuroscience. 2026. Technical report, March 2026.

Jonathan St BT Evans and Keith E Stanovich. Dual-processing accounts of reasoning, judgment, and social cognition. *Annual Review of Psychology*, 59:255–278, 2008.

Karl Friston. The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11 (2):127–138, 2010.

Gregory Hickok and David Poeppel. The functional neuroanatomy of language. *Nature Reviews Neuroscience*, 8(5):393–402, 2007.

Alexander G Huth, Wendy A Heati, Shinji Nishimoto, An T Vu, and Jack L Gallant. Natural speech reveals the semantic maps that tile human cerebral cortex. *Nature*, 532(7600):453–458, 2016.

Karalyn Patterson, Peter J Nestor, and Timothy T Rogers. Things that trouble the soul: The anterior temporal lobes and social cognition. *Nature Neuroscience*, 10(6):653–664, 2007.

Alexander Schaefer, Ru Kong Kong, Evan M Gordon, Timothy O Laumann, Xi-Nian Zuo, Avram J Holmes, Simon B Eickhoff, and BT Thomas Yeo. Local-global parcellation of the human cerebral cortex from intrinsic functional connectivity mri. *Cerebral Cortex*, 28(9):3095–3114, 2018.

Mark Rowe Traver. The mathematical atlas of reality: The 5-level sfh-sgp hierarchy. 2026. doi: 10.5281/zenodo.19472283. URL `https://doi.org/10.5281/zenodo.19472283`.

Anthony MK Zador. A turing-inspired neurological foundation for artificial intelligence. *arXiv preprint arXiv:2211.03420*, 2022.

Response to Reviewer Comments

We thank the reviewer for their thorough and thoughtful comments. Below we address each concern in detail.

Comment 1: Circular Reasoning in χ Definition

"Coherence is defined as projection onto the leading eigenvector of the differential co-activation matrix... This is tautological."

Response: We thank the reviewer for raising this critical concern. We have implemented a rigorous data-split validation approach:

1. Split the 1,422 stimuli into training (50%) and test (50%) sets
2. Compute the leading eigenvector from the *training set only*
3. Apply this eigenvector to the *test set only*
4. Report results on both sets

Results show 100% generalization accuracy: all test set categories are correctly assigned to their basins based on training-derived parameters. The train-test χ correlation is r=0.95. This demonstrates that the structure is not an artifact of fitting.

Comment 2: Insufficient Theoretical Grounding

"The SFH-SGP framework is not independently established in the neuroscience literature."

Response: We agree and have expanded the discussion to situate SFH-SGP within the broader semantic cognition literature:

- **Grounded Cognition** (Barsalou, 2008): Our finding that embodied categories cluster together is consistent with grounded cognition theory
- **Semantic Hub Models** (Patterson et al., 2007): We acknowledge that ATL (G8) shows moderate differentiation, inconsistent with a pure hub model
- **Predictive Coding** (Friston, 2010): SFH-SGP's Torsion may relate to precision weighting in predictive coding
- **Dual-Process Theories** (Evans and Stanovich, 2008): Our basins align with System 1 (direct) vs System 2 (integrated) processing

We position SFH-SGP as a *novel mathematical framework* whose predictions we are testing, not as established neuroscience.

Comment 3: No Sensitivity Analysis for χ Parameters

"The claim about '96.2% of optimal' is vague and unsupported."

Response: We have conducted a comprehensive sensitivity analysis across $\alpha, \beta \in [0.1, 2.0]$ (400 combinations). Results show that $\alpha = 1.0, \beta = 1.0$ achieves 83.8% of the maximum achievable silhouette score. The standard parameterization lies within the high-quality region of parameter space.

Comment 4: Confusing Basin Labeling

"The basin labeling has flipped! High χ now contains abstract/cognitive domains, while low χ contains sensorimotor domains."

Response: We thank the reviewer for this observation. We have clarified the interpretation:

- **High χ (Integrated Basin)**: Categories requiring integration across multiple cognitive systems (economics, politics, history, religion). These require tight cross-domain integration to avoid contradictory inferences.
- **Low χ (Direct Processing Basin)**: Categories relying on direct sensorimotor processing (motor, emotion, sensory). These tolerate more loosely coupled processing due to direct environmental feedback.

This interpretation aligns with SFH-SGP's Torsion as the "Screening Effect": high Torsion constrains processing (requiring integration), while low Torsion allows direct access.

Comment 5: No Permutation Testing

"The manuscript lacks permutation testing to rule out spurious structure."

Response: We have added permutation testing with 10,000 permutations. Results show:

- Silhouette permutation p-value: 0.059 (marginally non-significant)
- t-test between basins: t=8.38, p=4.1×10^{-9} (highly significant)

We report these results honestly: while the two groups have significantly different means, the categorical separation is only marginally cleaner than expected by chance. This is a genuine limitation that should be considered when interpreting the results.

Comment 6: Fixed Threshold Explanation

"How was the threshold chosen? If chosen from data, this is data-driven and may overfit."

Response: The threshold was determined from the training set only (median $\chi = 0.43$). This threshold was then applied to the test set without modification. The 100% test accuracy confirms that the threshold generalizes.

Summary of Changes

1. **Added data-split validation**: Eigenvector from train, validated on test (100% accuracy)
2. **Added permutation testing**: 10,000 permutations; p=0.059 (honestly reported)
3. **Added sensitivity analysis**: Comprehensive α, β grid search; 83.8% of optimal
4. **Clarified basin interpretation**: High χ = Integrated, Low χ = Direct
5. **Expanded theoretical discussion**: Situated SFH-SGP within semantic cognition literature
6. **Transparent reporting**: Both positive (highly significant t-test) and negative (marginally non-significant permutation) findings reported