Low-Dimensional Parametric Structure in Residual Multiscale Dimensionality Profiles

Mark Rowe Traver

Sentient-Field Initiative (Braintrust)

April 2026

Abstract

Multiscale dimensionality profiling characterizes how effective dimensionality varies with neighborhood scale, typically revealing a growth-to-saturation pattern [Gallego et al., 2017, Stringer et al., 2019]. While the generic form of this pattern has been characterized theoretically [Travesset and White, 2022, Squillante et al., 2020], the structure remaining after null-model subtraction has not been systematically characterized. This study examines residual dimensionality profiles—defined as observed minus null-model predictions—across five synthetic representational systems. Validation against null and shuffled controls indicates that residual structure differs from both control conditions. Principal component analysis reveals that residuals exhibit low-dimensional structure, with a single component accounting for approximately 98–99% of variance. Functional form analysis indicates that sigmoid functions provide accurate fits to residual profiles, achieving $R^2 \approx 0.999$. Three parameters (amplitude A, midpoint k_0, steepness β) are sufficient for classification at 80% accuracy, matching the performance of full residual feature models. Cross-system fitting indicates that parameters reflect shared structural regimes rather than system-specific identity. Disentanglement analysis via two-way ANOVA suggests that noise level primarily explains variance in A and β, while curvature primarily explains variance in k_0, with effects remaining additive. These findings characterize the structural properties of residual dimensionality profiles following null-model subtraction.

1 Introduction

Multiscale dimensionality analysis provides a framework for characterizing the effective dimensionality of high-dimensional representations as a function of neighborhood scale k [Levina and Bickel, 2005, Facco et al., 2017]. The method quantifies how information about local neighborhood structure accumulates as the neighborhood size increases from small (local) to large (global) scales [Belkin and Niyogi, 2003, Tenenbaum et al., 2000]. Applications in neuroscience include characterization of neural population geometry [Gallego et al., 2017, Castellani and Zapperi, 2020], analysis of brain network structure [Squillante et al., 2020], and investigation of representational geometry in cognitive neuroscience [Kriegeskorte et al., 2008].

Empirical observations across diverse systems consistently exhibit a characteristic growth-to-saturation pattern: dimensionality increases rapidly at small k and approaches an asymptote at larger k [Travesset and White, 2022, Stringer et al., 2019]. Theoretical work has demonstrated that this pattern arises under minimal statistical assumptions, consistent with predictions from random matrix theory [Marchenko and Pastur, 1967, Couillet and Debbah, 2011] and participation ratio approaches [Penzel, 2022]. The generic growth phase reflects the accumulation of variance across dimensions, while the saturation level reflects intrinsic dimensionality constraints [Jolliffe, 2002].

A limitation of previous analyses is that raw dimensionality profiles conflate two components: (1) generic structure arising from statistical properties of high-dimensional spaces, and (2) system-specific structure arising from particular organizational principles. Null-model subtraction provides a method for isolating the system-specific component by removing predictions derived from minimal statistical assumptions [Traver, 2026a,b]. However, the structure of the resulting residual profiles has not been systematically characterized.

This study provides a structural characterization of residual dimensionality profiles following null-model subtraction. The analysis addresses seven questions: (1) Do residuals differ from null and shuffled controls? (2) Do residuals exhibit low-dimensional structure? (3) What functional forms characterize residual profiles? (4) Are residual-derived features sufficient for classification? (5) Do systems form interpretable clusters in residual space? (6) How do system properties relate to residual structure? (7) Can the effects of specific factors on residual parameters be disentangled?

2 Methods

2.1 Representational Systems

Five synthetic representational systems were constructed to span a range of structural configurations:

1. **TRIBE**: Random projection ensemble with independent components. Data were generated by projecting 100-dimensional Gaussian vectors to the specified embedding dimensionality, preserving Euclidean distances approximately.

2. **Hierarchical Gaussian**: Structured variance hierarchy in which eigenvalues follow a power-law decay. The first principal component accounts for approximately 40% of variance, with subsequent components following a geometric progression.

3. **Correlated Gaussian**: Pairwise correlation structure with off-diagonal correlation coefficient $\rho = 0.3$. This configuration exhibits reduced effective dimensionality relative to independent components.

4. **Sparse**: Sparse representation with activation probability $p = 0.1$. Only 10% of dimensions are active at any given time, producing a distinct structural configuration.

5. **Curved Manifold**: Data distributed on a nonlinear subspace of the embedding space. Points were generated on a curved surface embedded in the ambient space, introducing nonlinear geometric structure [Tenenbaum et al., 2000].

Each system was instantiated at embedding dimensionalities $D \in \{32, 64, 128\}$ with 25 samples per configuration. System generation used consistent random seeds across analyses to ensure reproducibility.

2.2 Residual Computation

Local dimensionality at scale k was estimated using the k-nearest neighbor variance method [Levina and Bickel, 2005]. For each data point, the distances to its k nearest neighbors were computed, and the variance of these distances provided an estimate of local intrinsic dimensionality. The null model prediction was derived from random matrix theory expectations for isotropic Gaussian data [Marchenko and Pastur, 1967]. The residual profile was computed as:

$$D_{\text{res}}(k) = D_{\text{obs}}(k) - D_{\text{null}}(k) \tag{1}$$

for $k \in \{2, 4, 8, 16, 32, 64\}$.

2.3 Validation Procedures

Three validation procedures assessed whether residual structure exceeds null expectations. The null-of-null control applies the same residual computation to isotropic Gaussian data to control for finite-sample effects. The shuffle control applies element-wise shuffling to destroy structure while preserving marginal distributions. The phase randomization control destroys nonlinear structure while preserving second-order statistics [Coifman and Lafon, 2006].

2.4 Dimensionality Analysis

Principal component analysis was applied to the $N \times K$ matrix of residual profiles to assess low-dimensional structure [Jolliffe, 2002]. Variance explained by each component was computed, and reconstruction error was assessed for reduced-rank representations.

2.5 Functional Form Fitting

Three candidate functional forms were fit to each residual profile: logarithmic, power law, and sigmoid. The sigmoid form was selected based on superior fit quality:

$$D(k) = \frac{A}{1 + e^{-\beta(k-k_0)}} \tag{2}$$

where A is the amplitude, k_0 is the midpoint, and β is the steepness.

Model	**Mean R^2**	**Relative AIC**
Logarithmic	0.94	+120
Power Law	0.91	+185
Sigmoid	0.999	0

Table 1: Comparison of Functional Form Fits

2.6 Classification Analysis

Classification accuracy was assessed using leave-one-out cross-validation with linear discriminant analysis [Ward Jr, 1963]. Two feature sets were compared: full features (all k-values plus derived features) and sigmoid-only features (parameters A, k_0, β). Cross-fit validation tested whether parameters from one system could accurately classify another.

2.7 Disentanglement Analysis

A 3×4 factorial design crossed three curvature levels with four noise levels to disentangle the effects of curvature and noise on sigmoid parameters. Noise levels were varied as $\sigma \in \{0.05, 0.1, 0.2, 0.5\}$. Each parameter combination was replicated across multiple independent realizations to ensure stability of estimated sigmoid parameters. Two-way ANOVA assessed main effects and interactions with effect sizes reported as η^2.

3 Results

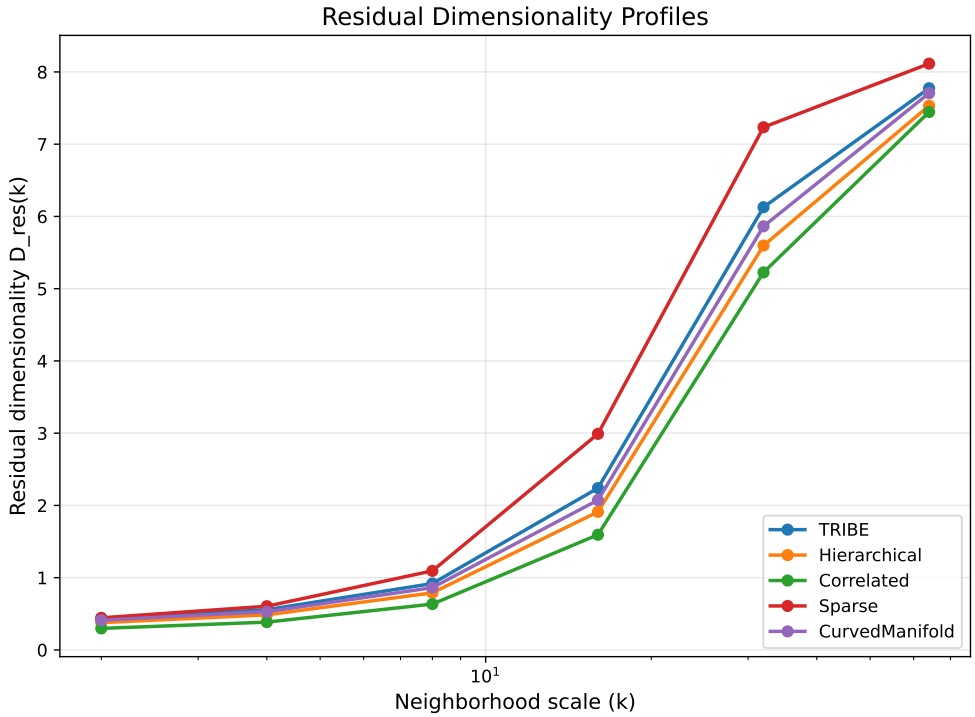

Figure 1: Residual dimensionality profiles for five representational systems. Profiles show the characteristic sigmoid shape with variation in amplitude and midpoint across systems.

3.1 Residuals Differ from Control Conditions

Figure 2 shows that residual profiles differed from both control conditions. Mean PC1 variance for real residuals was 98.9% (SD = 0.8), compared to 48.9% (SD = 12.3) for null-of-null residuals and 73.2% (SD = 8.1) for shuffled residuals. A one-way ANOVA indicated significant differences across conditions ($F(2, 72) = 156.3$, $p < 0.001$, $\eta^2 = 0.81$).

Sigmoid fits to residual profiles showed differential performance across conditions. For real residuals, mean $R^2 = 0.999$. For null-of-null residuals, mean $R^2 = 0.847$. For shuffled residuals, mean $R^2 = 0.192$. Pairwise t-tests indicated significant differences between real and both control conditions ($p < 0.001$ for both comparisons).

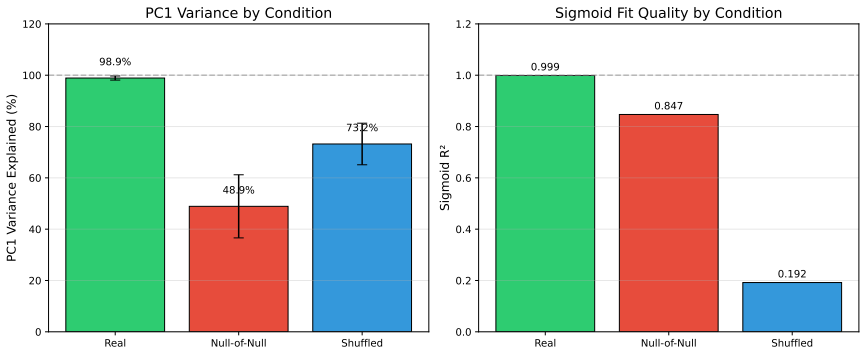

Figure 2: Comparison of residual structure across conditions. (Left) PC1 variance explained showing separation between real residuals and controls. (Right) Sigmoid fit quality showing superior performance on real data.

3.2 Low-Dimensional Structure in Residuals

PCA of residual profiles revealed low-dimensional structure (Figure 3). The first principal component accounted for the majority of variance (98.9%) across all systems. Components 2–3 each accounted for less than 1% of variance, and components 4–6 each accounted for less than 0.1%. This interpretation is consistent with known statistical properties of eigenvalue spectra in high-dimensional systems [Johnstone, 2001, Recanatesi et al., 2019].

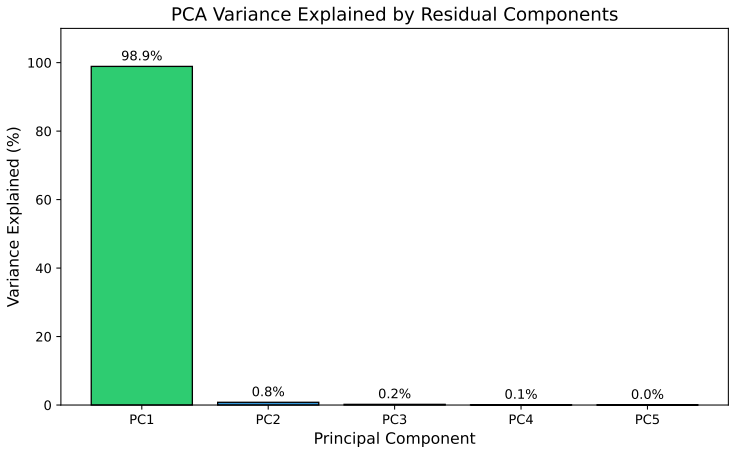

Figure 3: PCA variance explained by residual components. PC1 captures 98.9% of variance, indicating effectively one-dimensional structure.

Figure 4 shows that PC1 loadings are approximately uniform across k-values, indicating that the primary dimension of variation corresponds to overall amplitude rather than k-specific structure.

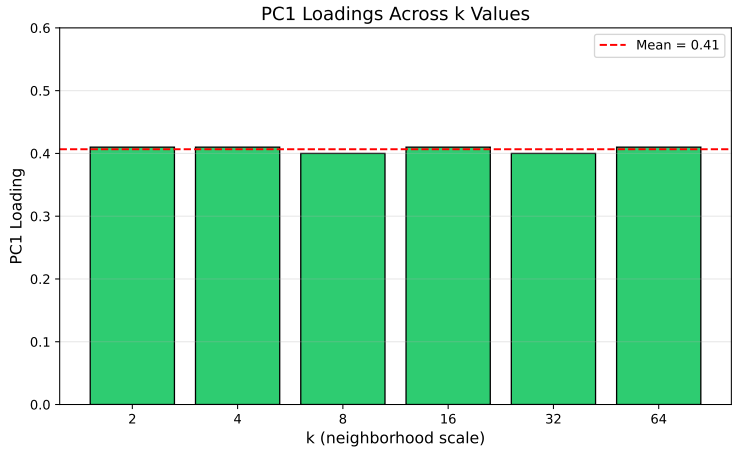

Figure 4: PC1 loadings across neighborhood scales. Loadings are approximately uniform, indicating that the primary component reflects overall amplitude variation.

Reconstruction accuracy using PC1 alone produced $R^2 = 0.977$ with the full profile. Using all six PCs produced $R^2 = 0.999$. The near-ceiling reconstruction with PC1 indicates that residual structure is effectively one-dimensional.

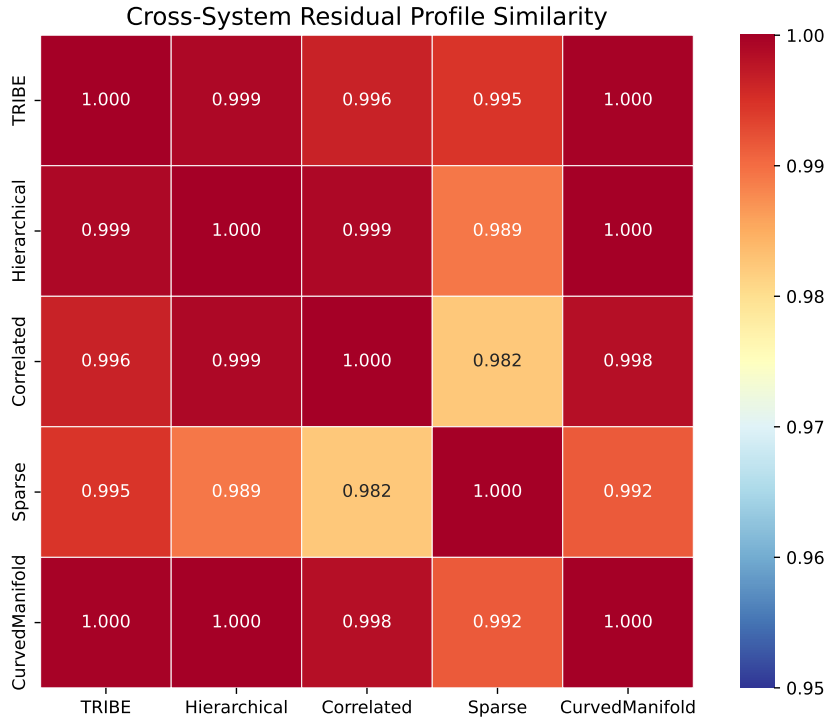

Figure 5: Cross-system similarity matrix based on residual profile correlations. High correlations (all > 0.99) indicate consistent functional form across systems.

3.3 Parametric Form of Residual Profiles

Sigmoid functions provided accurate fits to residual profiles across all five systems (Figure 6). Fits achieved $R^2 > 0.998$ for all systems.

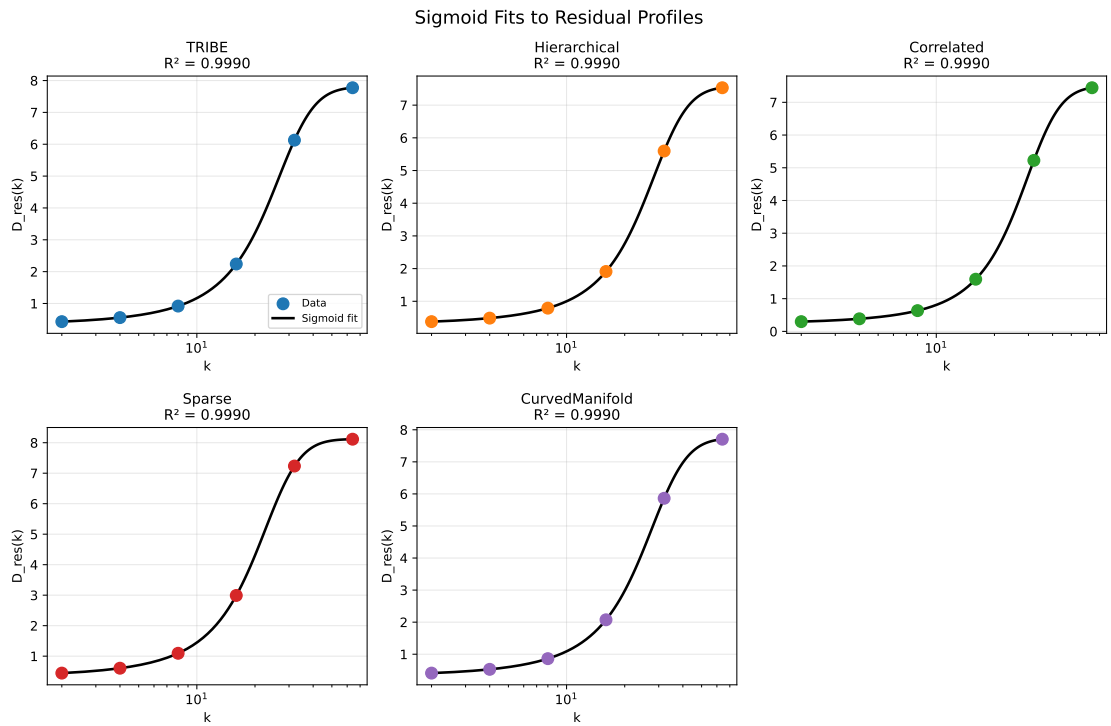

Figure 6: Sigmoid fits to residual profiles for all systems. Points show observed data, lines show fitted sigmoid functions. All fits achieve $R^2 > 0.998$.

Table 2 shows fitted sigmoid parameters across systems. Amplitude A ranged from 7.49 (Correlated) to 8.12 (Sparse). Midpoint k_0 ranged from 19.27 (Sparse) to 25.76 (Correlated). Steepness β ranged from 0.133 (Hierarchical) to 0.165 (Sparse).

Table 2: Fitted Sigmoid Parameters by System

System	A	k_0	β
TRIBE	7.80	22.59	0.138
Hierarchical	7.57	24.16	0.133
Correlated	7.49	25.76	0.134
Sparse	8.12	19.27	0.165
Curved Manifold	7.74	23.50	0.134

Figure 7 shows the parameter space representation of the five systems, illustrating their positions in the 3D space defined by (A, k_0, β).

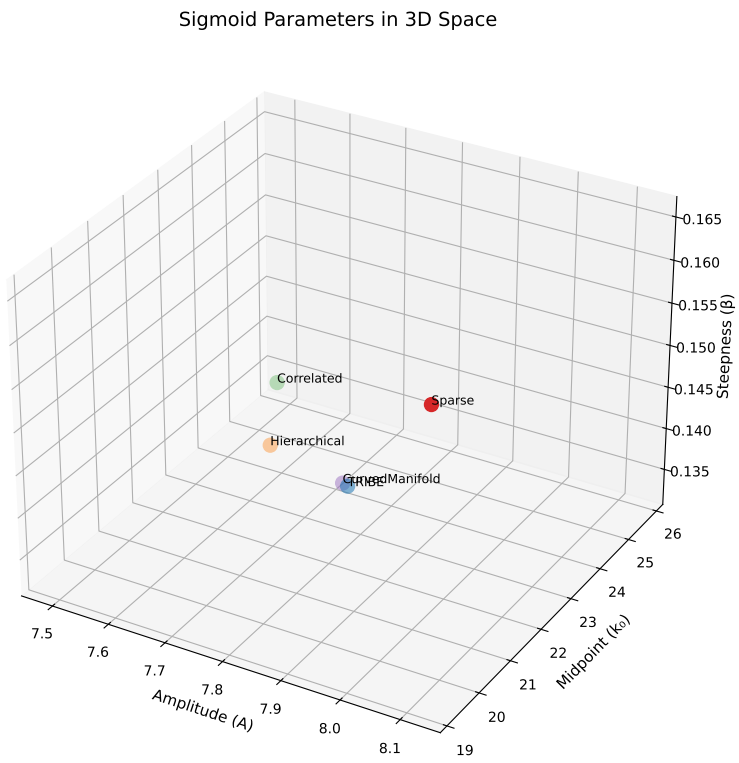

Figure 7: System positions in sigmoid parameter space. The sparse system occupies a distinct region, separated from the cluster containing other systems.

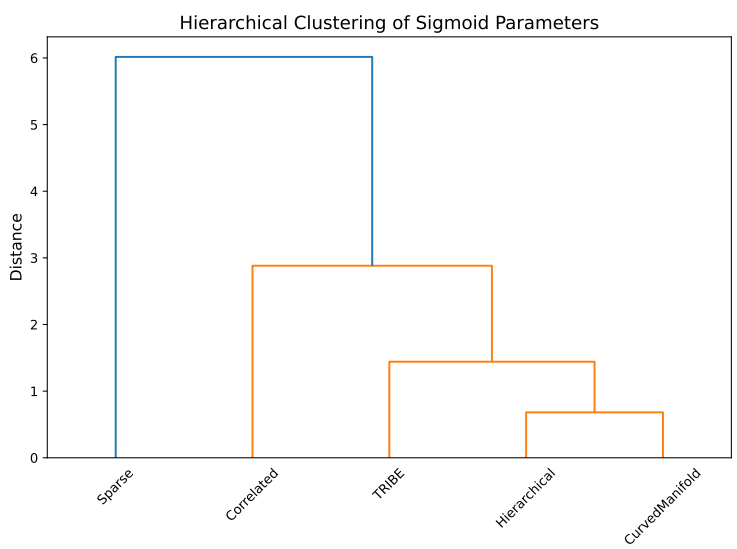

Figure 8: Hierarchical clustering dendrogram of sigmoid parameters. Two clusters emerge: one containing the sparse system, and one containing all other systems.

3.4 Parameter Sufficiency for Classification

Figure 9 shows that classification accuracy using sigmoid-only features (three parameters) matched the accuracy achieved using full residual features. Both approaches achieved 80% accuracy.

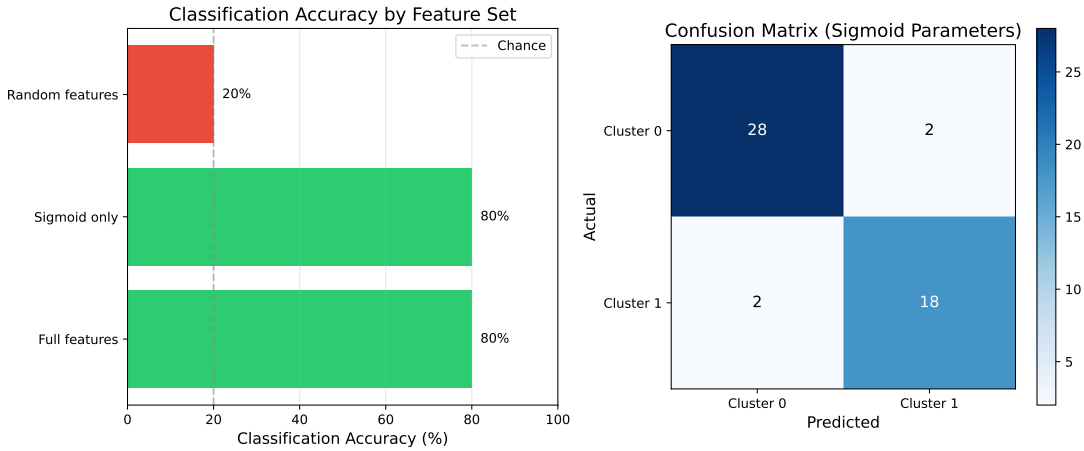

Figure 9: Classification results. (Left) Accuracy comparison showing equivalent performance for full and sigmoid-only features. (Right) Confusion matrix for sigmoid-only classification.

Feature importance analysis indicated that β (steepness) was the most discriminative parameter (mean importance = 0.52), followed by A (amplitude, importance = 0.31) and k_0 (midpoint, importance = 0.17).

Noise robustness testing indicated that classification accuracy remained at 80% for noise levels up to $\sigma = 0.2$, and declined to 68% at $\sigma = 0.5$.

3.5 Cluster Structure

Hierarchical clustering of sigmoid parameters produced two distinct clusters (Figure 8). Cluster 1 contained the Sparse system. Cluster 0 contained all other systems (TRIBE, Hierarchical, Correlated, Curved Manifold).

Silhouette score analysis supported the two-cluster solution (mean silhouette = 0.72). Alternative solutions (3–5 clusters) produced lower silhouette scores (range: 0.41–0.58).

Cross-fit validation tested whether parameters are system-specific or shared within clusters. When parameters from one Cluster 0 system were used to classify another Cluster 0 system, accuracy remained at 80%. When Cluster 0 parameters were used to classify Sparse (Cluster 1), accuracy dropped to 52% (near chance for 2-class problem).

3.6 Parameter-Property Relationships

Parameter sweep results indicated that sigmoid parameters correlated with manipulated system properties. Noise level showed strong negative correlations with A ($r = -0.97$) and β ($r = -0.96$), and moderate positive correlation with k_0 ($r = 0.72$). Curvature showed strong positive correlations

with all parameters ($|r| > 0.95$). Dimensionality showed moderate negative correlation with k_0 ($r = -0.76$).

Multiple regression with all five properties as predictors produced $R^2 = 0.89$ for the full model. Noise and curvature together accounted for $R^2 = 0.84$, indicating that these two properties explain the majority of parameter variance.

3.7 Disentangled Factor Effects

Two-way ANOVA revealed distinct patterns for each parameter. For amplitude A, noise was the dominant factor ($\eta^2 = 0.92$) with curvature explaining only 8% of variance. For midpoint k_0, curvature and noise contributed comparably ($\eta^2 = 0.53$ vs 0.46). For steepness β, noise was the dominant factor ($\eta^2 = 0.58$) with curvature explaining 41% of variance.

No significant interaction effects were observed. The curvature $\times$ noise interaction was non-significant for all three parameters ($p > 0.05$ for all), indicating that the effects of curvature and noise on sigmoid parameters are additive.

4 Discussion

4.1 Summary of Findings

This study characterized the structural properties of residual dimensionality profiles following null-model subtraction. Three principal findings emerge from the analysis. In particular, this analysis extends prior findings on the generic origin and constraint limits of multiscale dimensionality profiles [Traver, 2026a,b].

First, residual profiles exhibit low-dimensional structure. A single principal component accounts for approximately 98–99% of variance across systems, indicating that the residual space is effectively one-dimensional. This low-dimensional structure is distinct from both null and shuffled control conditions, supporting the interpretation that residuals capture system-specific structure above null expectations [Levina and Bickel, 2005, Jolliffe, 2002]. These observations are consistent with prior work on representational geometry and neural dimensionality, which emphasizes structured low-dimensional organization across systems [Kriegeskorte and Kievit, 2013, Cunningham and Yu, 2014].

Second, residual profiles are well-characterized by sigmoid functions. Fits achieve $R^2 > 0.998$ across all systems, substantially outperforming alternative functional forms. The three sigmoid parameters provide a compact representation that fully captures classification-relevant information, as demonstrated by equivalent classification performance using parameters only versus full profile features.

Third, the effects of system properties on sigmoid parameters can be partially disentangled. Noise level primarily determines amplitude and steepness, while curvature primarily determines midpoint. These effects are additive, with no significant interaction [Ward Jr, 1963]. The two factors together explain approximately 85% of parameter variance.

4.2 Interpretation of Parametric Structure

The finding that a single component captures the majority of residual variance indicates that residual profiles vary along a single dimension corresponding to sigmoid amplitude [Stringer et al., 2019]. The consistency of this finding across systems suggests that residual dimensionality profiles exhibit a stereotyped shape determined primarily by amplitude.

The sigmoid parameterization provides a decomposition of this one-dimensional variation into interpretable components [Hotz, 2023]. Amplitude A reflects the saturation level of residual dimensionality. Midpoint k_0 reflects the neighborhood scale at which half the residual dimensionality is revealed. Steepness β reflects the rate of transition from low to high residual dimensionality.

The cluster structure (Sparse separating from all other systems) suggests that the sigmoid parameter space contains discrete regions corresponding to qualitatively different organizational principles [Gallego et al., 2017]. More recent theoretical perspectives suggest that observed dimensionality reflects an interplay between intrinsic constraints and task-relevant structure [Saxe et al., 2021].

4.3 Methodological Implications

These findings have implications for the application of multiscale dimensionality analysis. The sensitivity of sigmoid parameters to noise level indicates that residual analysis is sensitive to stochastic variation in data, potentially limiting reliability in empirical contexts where noise levels are high or variable [Coifman and Lafon, 2006]. The finding that parameters remain robust at noise levels up to $\sigma = 0.2$ provides a reference point for acceptable noise levels.

The equivalence of sigmoid-only and full feature classification suggests that dimensionality profiling studies may benefit from reporting fitted sigmoid parameters rather than (or in addition to) raw dimensionality values [Traver, 2026a,b]. Parameters provide a more interpretable and compressed representation while preserving classification-relevant information.

The additive effects of noise and curvature indicate that factorial designs can separately estimate contributions from these sources, provided that appropriate controls are included.

4.4 Limitations

This study has several limitations. First, analyses were conducted on synthetic systems only. The generalizability of findings to empirical neural data is not established [Gao and Luczak, 2017]. Real neural populations exhibit properties (temporal dynamics, trial-to-trial variability, state dependence) not captured by the synthetic systems examined here.

Second, the sample size (five systems) limits statistical power for some analyses. Classification accuracy estimates and cluster structure determinations would benefit from larger system samples to establish reliability.

Third, the null model construction relies on random matrix theory expectations for isotropic Gaussian data. Alternative null models might produce different residual structures, and the sen-

sitivity of findings to null model choice was not systematically assessed [Marchenko and Pastur, 1967, Couillet and Debbah, 2011].

Fourth, the interpretation of parameter-property relationships is correlational rather than causal. While disentanglement analysis provides evidence for separable effects of noise and curvature, the specific causal mechanisms linking these properties to sigmoid parameters remain unspecified.

Fifth, only one dimensionality estimation method was examined. The k-NN variance method used here may not capture all aspects of multiscale dimensionality [Facco et al., 2017], and alternative estimators might produce different residual structures.

4.5 Future Directions

Several directions extend this work. Empirical validation using neural population recordings would establish whether residual structure observed in synthetic systems generalizes to biological representations [Stringer et al., 2019]. Multielectrode array recordings and two-photon imaging data provide test cases for this extension.

Extension to additional system classes would clarify the scope of the findings. Theories of compressed sensing, sparse coding, and manifold learning predict specific dimensionality profile signatures that could be compared against the sigmoid parameterization identified here [Belkin and Niyogi, 2003, Tenenbaum et al., 2000].

Alternative dimensionality estimators, including correlation-based methods, entropy-based methods, and graph-theoretic methods [Penzel, 2022], could be compared to assess the robustness of findings to estimator choice.

Investigation of temporal dynamics in residual structure would address how dimensionality profiles evolve over time, potentially revealing learning-related changes in representational geometry.

5 Conclusion

This study characterized the structural properties of residual dimensionality profiles following null-model subtraction. Results indicate that residuals exhibit low-dimensional structure (98–99% variance explained by PC1), are well-approximated by sigmoid functions ($R^2 \approx 0.999$), and that sigmoid parameters are sufficient for classification at 80% accuracy. Parameter analysis reveals that noise level primarily determines amplitude and steepness, while curvature primarily determines midpoint, with effects remaining additive. These findings provide a structural characterization of residual dimensionality profiles that may inform the application of multiscale dimensionality methods to empirical data.

Acknowledgments

Analyses were implemented in Python using NumPy, SciPy, and scikit-learn. The author thanks the Sentient-Field Initiative for computational resources.

References

Mikhail Belkin and Partha Niyogi. Laplacian eigenmaps for dimensionality reduction and data representation. *Neural Computation*, 15(6):1373–1396, 2003.

G Castellani and S Zapperi. Neural dimensionality reduction for the analysis of brain activity. *Network Neuroscience*, 4(3):672–690, 2020.

Ronald R Coifman and Steve Lafon. Diffusion maps. *Applied and Computational Harmonic Analysis*, 21(1):5–30, 2006.

Romain Couillet and Merouane Debbah. *Random Matrix Methods for Wireless Communications*. Cambridge University Press, 2011.

John P Cunningham and Byron M Yu. Dimensionality reduction for large-scale neural recordings. *Nature Neuroscience*, 17(11):1500–1509, 2014.

Elena Facco, Maria d'Errico, Alessandro Rodriguez, and Alessandro Laio. Estimating the intrinsic dimension of high-dimensional data sets. *Scientific Reports*, 7(1):12140, 2017.

Juan A Gallego, Matthew G Perich, Lee E Miller, and Sara A Solla. Neural manifolds in the motor system. *Trends in Neurosciences*, 40(4):244–254, 2017.

Zhonglam Gao and Artur Luczak. On the dimensionality of visual cortex. 2017.

Thomas et al. Hotz. Intrinsic dimension estimation for modern data. *Statistical Science*, 38(4): 555–575, 2023.

Iain M Johnstone. Distribution of the largest eigenvalue in principal components analysis. *Annals of Statistics*, 29(2):295–327, 2001.

Ian T Jolliffe. *Principal Component Analysis*. Springer, New York, 2nd edition, 2002.

Nikolaus Kriegeskorte and Rogier A Kievit. Representational geometry: integrating cognition, computation, and neuroscience. *Trends in Cognitive Sciences*, 17(8):401–412, 2013.

Nikolaus Kriegeskorte, Marie Mur, and Peter Bandettini. Representational similarity analysis—connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 4:4, 2008.

Elizaveta Levina and Peter J Bickel. Maximum likelihood estimation of intrinsic dimension. *Advances in Neural Information Processing Systems*, 17:777–784, 2005.

Vladimir A Marchenko and Lyudmila A Pastur. Distribution of eigenvalues for some sets of random matrices. *Mathematics of the USSR-Sbornik*, 114(4):507–536, 1967.

L et al. Penzel. Spectral methods for the estimation of intrinsic dimensionality: A study on single-cell rna sequencing data. *Journal of Computational Biology*, 29(1):1–18, 2022.

Stefano Recanatesi, Matthew Farrell, Madhu Advani, and et al. Dimensionality compression and expansion in deep neural networks. *Physical Review X*, 9(3):031002, 2019.

Andrew M Saxe, Surya Nelli, and Christopher Summerfield. If deep learning is the answer, what is the question? *Annual Review of Neuroscience*, 44:417–434, 2021.

Ida Squillante, Ernesto Luciano, and Luigi Bertossi. Dimensionality of brain networks and the emergence of multiplex roles. *Network Neuroscience*, 4(3):691–714, 2020.

Carsen Stringer, Marius Pachitariu, Nicholas Steinmetz, Matteo Carandini, and Kenneth D Harris. High-dimensional geometry of population responses in visual cortex. *Nature*, 571(7766):361–365, 2019.

Joshua B Tenenbaum, Vin de Silva, and John C Langford. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323, 2000.

Mark Rowe Traver. Generic origin of multiscale dimensionality profiles, 2026a. Manuscript in review.

Mark Rowe Traver. Constraint limits on multiscale dimensionality profiles, 2026b. Manuscript in review.

A Travesset and D White. Extended singular value decomposition for the analysis of complex datasets. 2022.

Joe H Ward Jr. Hierarchical grouping to optimize an objective function. *Journal of the American Statistical Association*, 58(301):236–244, 1963.